



CAMBRIDGE

Lecture 4: Neural Networks

Stefan Bucher

MACHINE LEARNING IN ECONOMICS
UNIVERSITY OF CAMBRIDGE

Stefan Bucher

 Open in Colab

Prince (2023, chaps. 3, 4, 7).¹

- . Figures taken or adapted from Prince (2023). All rights belong to the original author and publisher. These materials are intended solely for educational purposes.

Shallow Neural Networks



Prince (2023, chap. 3)

Example


$$f[x, \phi] = \phi_0 + \underbrace{\phi_1 a(\theta_{10} + \theta_{11}x)}_{h_1} + \underbrace{\phi_2 a(\theta_{20} + \theta_{21}x)}_{h_2} + \underbrace{\phi_3 a(\theta_{30} + \theta_{31}x)}_{h_3}$$

$$a(z) = \text{ReLU}(z) = \max\{z, 0\}$$

Neural networks (shallow & deep!) with ReLU represent continuous piecewise linear functions.




Activation Patterns




Joints where one hidden unit becomes (in)active.

Multivariate Output

a)




b)



Shallow Neural Networks

Map n_{in} -dimensional input to n_{out} -dimensional output.


$$y_j = \phi_{j0} + \sum_{d=1}^D \phi_{jd} \underbrace{a\left(\theta_{d0} + \sum_{i=1}^{n_{in}} \theta_{di} x_i\right)}_{h_d}, \quad j = 1, \dots, n_{out}$$



Shallow Networks (3.3 & 3.8)

Universal Approximation Thm

$\exists$ shallow neural network with a single layer of sufficiently many hidden units that can approximate, with arbitrary precision, any continuous function from a compact subset of $\mathbb{R}^{n_{in}}$ to $\mathbb{R}^{n_{out}}$.




Deep Neural Networks

Prince (2023, chap. 4)

Why Depth? (4.1)


- Universal approximation theorem...
- Deep: more linear regions for same # parameters

Extra Layer is Folding




Extra Layer is Clipping (4.5)

Composing as Special Case




Composing in 2D


a)




b)



c)



d)




Deep Neural Network



$$\mathbf{h}_1 = \mathbf{a}(\beta_0 + \mathbf{\Omega}_0 \mathbf{x})$$

$$\mathbf{h}_k = \mathbf{a}(\beta_{k-1} + \mathbf{\Omega}_{k-1} \mathbf{h}_{k-1}), \quad k = 2, \dots, K$$

$$\mathbf{y} = \beta_K + \mathbf{\Omega}_K \mathbf{h}_K$$



Depth Efficiency



Backpropagation


Prince (2023, chaps. 7.1–7.4)

Computing Gradients with Chain Rule

Stochastic Gradient Descent

$$\phi_{t+1} \leftarrow \phi_t - \alpha \sum_{i \in B_t} \nabla_{\phi} l_i(\phi_t)$$

requires computing $\frac{\partial l_i}{\partial \beta_k}$ and $\frac{\partial l_i}{\partial \Omega_k}$ at each layer $k \in \{0, \dots, K\}$.




$$\frac{\partial l_i}{\partial \Omega_k} = \frac{\partial l_i}{\partial \mathbf{f}_k} \frac{\partial \mathbf{f}_k}{\partial \Omega_k} = \frac{\partial \mathbf{h}_{k+1}}{\partial \mathbf{f}_k} \frac{\partial \mathbf{f}_{k+1}}{\partial \mathbf{h}_{k+1}} \frac{\partial l_i}{\partial \mathbf{f}_{k+1}} \mathbf{h}_k^T$$

where activation $\mathbf{h}_{k+1}$ results from pre-activation $\mathbf{f}_k$.

Forward Pass

Gradients of $l_i = l(\mathbf{f}[\mathbf{x}_i, \phi], \mathbf{y}_i)$ w.r.t. weights depend on activations $\mathbf{h}_k$, which we compute first.




$$\mathbf{h}_k = \mathbf{a}(\mathbf{f}_{k-1})$$

$$\mathbf{f}_k = \beta_k + \Omega_k \mathbf{h}_k$$

Backward Pass

Then compute gradients backwards using chain rule.



$$\frac{\partial l_i}{\partial \mathbf{f}_k} = \frac{\partial \mathbf{h}_{k+1}}{\partial \mathbf{f}_k} \frac{\partial \mathbf{f}_{k+1}}{\partial \mathbf{h}_{k+1}} \frac{\partial l_i}{\partial \mathbf{f}_{k+1}} = \mathbb{I}[\mathbf{f}_k > 0] \odot \boldsymbol{\Omega}_{k+1}^T \frac{\partial l_i}{\partial \mathbf{f}_{k+1}}$$

$$\frac{\partial l_i}{\partial \boldsymbol{\Omega}_k} = \frac{\partial \mathbf{f}_k}{\partial \boldsymbol{\Omega}_k} \frac{\partial l_i}{\partial \mathbf{f}_k} = \frac{\partial l_i}{\partial \mathbf{f}_k} \mathbf{h}_k^T$$

Backpropagation Algorithm

Forward pass computes activations

$$\begin{aligned}\mathbf{f}_0 &= \beta_0 + \mathbf{\Omega}_0 \mathbf{x}_i \\ \mathbf{h}_k &= \mathbf{a}(\mathbf{f}_{k-1}) & k \in \{1, \dots, K\} \\ \mathbf{f}_k &= \beta_k + \mathbf{\Omega}_k \mathbf{h}_k & k \in \{1, \dots, K\}\end{aligned}$$

Backward pass computes gradients


$$\begin{aligned}\frac{\partial l_i}{\partial \beta_k} &= \frac{\partial l_i}{\partial \mathbf{f}_k} & k \in \{K, K-1, \dots, 0\} \\ \frac{\partial l_i}{\partial \mathbf{\Omega}_k} &= \frac{\partial l_i}{\partial \mathbf{f}_k} \mathbf{h}_k^T & k \in \{K, K-1, \dots, 0\} \\ \frac{\partial l_i}{\partial \mathbf{f}_{k-1}} &= \mathbb{I}[\mathbf{f}_{k-1} > 0] \odot \left(\mathbf{\Omega}_k^T \frac{\partial l_i}{\partial \mathbf{f}_k} \right) & k \in \{K, K-1, \dots, 1\}\end{aligned}$$

where $l_i = l(\mathbf{f}[\mathbf{x}_i, \phi], \mathbf{y}_i)$ and $\mathbf{h}_0 = \mathbf{x}_i$ (i.e. do for each training example in batch).¹

- . This is done in parallel across a batch, so we get tensors instead of matrices.

PyTorch Implementation

Moons Dataset



Neural Network Implementation

```

1 import torch
2 import torch.nn as nn
3
4 class MyFirstNet(nn.Module):
5     def __init__(self):
6         super(MyFirstNet, self).__init__()
7         self.layers = nn.Sequential(
8             nn.Linear(2, 16),
9             nn.ReLU(),
10            nn.Linear(16, 2),
11        )
12
13    def forward(self, x):
14        return self.layers(x)
15
16    def predict(self, x):
17        output = self.forward(x)
18        return torch.argmax(output, 1)
19
20    def train(self, X, y):
21        loss_function = nn.CrossEntropyLoss()
22        optimizer = torch.optim.SGD(model.parameters(), lr=1e-2)
23
24        epochs = 15000

```

```

MyFirstNet(
  (layers): Sequential(
    (0): Linear(in_features=2, out_features=16, bias=True)
    (1): ReLU()
    (2): Linear(in_features=16, out_features=2, bias=True)
  )

```

Sample input:

```
tensor([[ 0.9610, -0.4991],
        [ 0.8780, -0.3601],
        [-1.0799,  0.1594],
        [ 0.8342,  0.3770],
        [ 0.5005, -0.5789]], device='mps:0')
```


Network output:

```
tensor([[ 0.1133, -0.0112],
        [ 0.1046, -0.0326],
        [ 0.3253, -0.3412],
        [ 0.0042,  0.0007],
```

Training

```
1 losses = model.train(X, y)
```

```
Epoch 0 loss is 0.6564172506332397  
Epoch 1000 loss is 0.28181493282318115  
Epoch 2000 loss is 0.2414126992225647  
Epoch 3000 loss is 0.2243497371673584  
Epoch 4000 loss is 0.21077106893062592  
Epoch 5000 loss is 0.197674959897995  
Epoch 6000 loss is 0.18398109078407288  
Epoch 7000 loss is 0.16948094964027405  
Epoch 8000 loss is 0.15426474809646606  
Epoch 9000 loss is 0.13909949362277985  
Epoch 10000 loss is 0.1247006431221962  
Epoch 11000 loss is 0.11129634082317352  
Epoch 12000 loss is 0.09911473095417023  
Epoch 13000 loss is 0.08828072249889374  
Epoch 14000 loss is 0.07882219552993774
```



Another Example

```
1 import torch, torch.nn as nn
2 from torch.utils.data import TensorDataset, DataLoader
3 from torch.optim.lr_scheduler import StepLR
4
5 # input size, hidden layer size, output size
6 D_i, D_k, D_o = 10, 40, 5
7 # create model with two hidden layers
8 model = nn.Sequential(
9     nn.Linear(D_i, D_k),
10    nn.ReLU(),
11    nn.Linear(D_k, D_k),
12    nn.ReLU(),
13    nn.Linear(D_k, D_o))
14
15 # He initialization of weights
16 def weights_init(layer_in):
17     if isinstance(layer_in, nn.Linear):
18         nn.init.kaiming_normal_(layer_in.weight)
19         layer_in.bias.data.fill_(0.0)
20 model.apply(weights_init)
21 # choose least squares loss function
22 criterion = nn.MSELoss()
23 # construct SGD optimizer and initialize learning rate and momentum
24 optimizer = torch.optim.SGD(model.parameters(), lr = 0.1, momentum=0.9)
```

Epoch 99, loss 4.419

References

Prince, Simon J. D. 2023. *Understanding Deep Learning*. Cambridge, Massachusetts: The MIT Press.